



BITS Pilani
Hyderabad Campus

Database Systems (CSF 212)

Prof.R.Gururaj
CS&IS Dept.

Indexing (Ch.17 of T1)



Content

- ☐ *What is Indexing*
- ☐ *Primary and Secondary indexes*
- ☐ *Dense and Sparse Indexing*
- ☐ *Multilevel Indexing*
- ☐ *Designing Primary and Multilevel Indexes*
- ☐ *What is Tree Indexing*
- ☐ *B+ tree*
- ☐ *Inserting and deleting keys into B+ Trees*
- ☐ *Constructing a B+ tree*
- ☐ *Designing a B+ Tree node structure*

Introduction to Indexing

Indexes (Access Structures) are used to speed up the retrieval of records in response to certain search conditions.

Ex: Get all employees details where age=50;
Get employee details for eid=329; etc.

In real world databases, indexes may be too large to be handled efficiently.

Hence some sophisticated techniques are to be used.

The criteria for evaluating the hashing or indexing techniques –

- ❖ Access time
- ❖ Insertion time (new indexes or new records)
- ❖ Deletion time
- ❖ space overhead

Some times more than one indexing may be required for a file.

The attribute /field used for constructing index structure for a file is called an '*indexing field/attribute*' .

If the index field is a key, it is called as search key or indexing key.

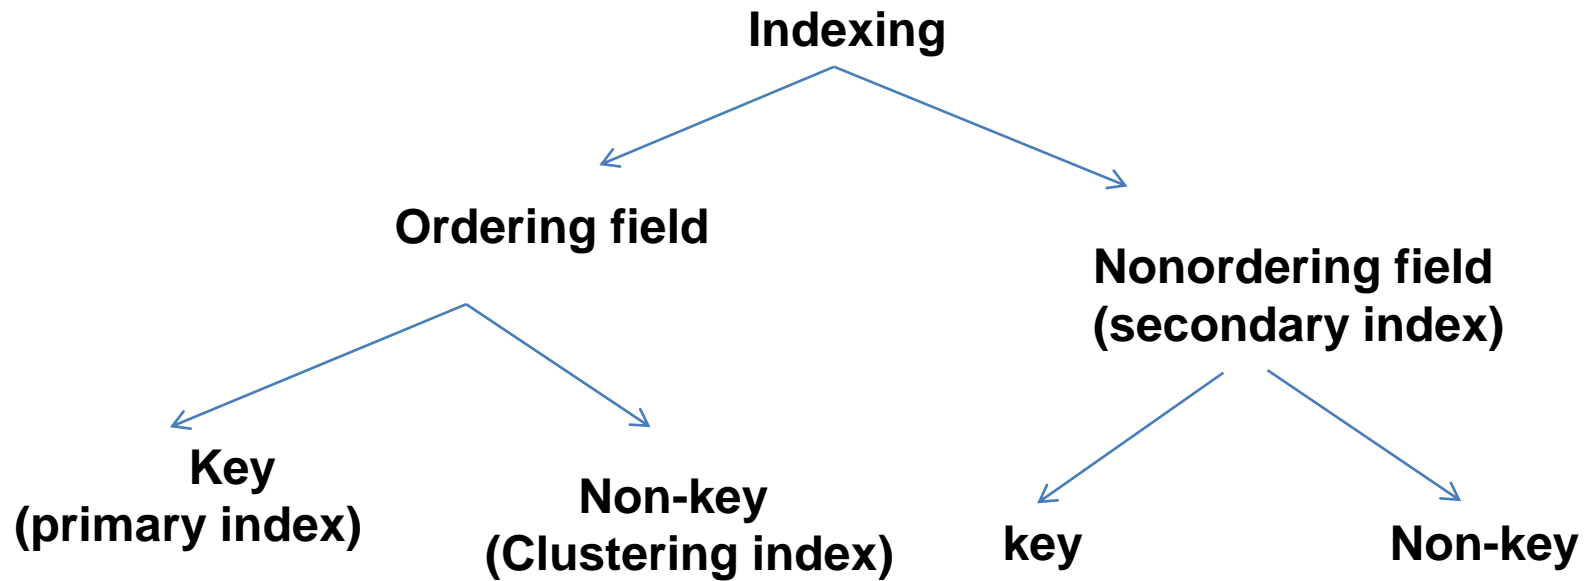
Indexes on key attributes:

1. Built on ordering key(PK) – Primary index
2. Non-ordering Key - Secondary index on key attribute

Indexes on non-key attributes:

1. Ordering non-key -- Clustering Index
2. Non-ordering non-key attribute – Secondary index on non-key

Hence, a file can have at most one primary index or one clustering index, but not both.



Data record: Similar kind of records(of a relation/table) are stored in a single file containing blocks. These are called data records and will have fields specified on the relation.

Index record: Like data records, index records are also stored in database. Any index record normally has two fields.

Value	Pointer
-------	---------

Key value

Location address of
the record containing
the key

Dense Index : In this, an index record appears for every data file record.

Sparse Index : Index records are created only for some data file records. This occupies less space.

A primary index and clustering index are non-dense.

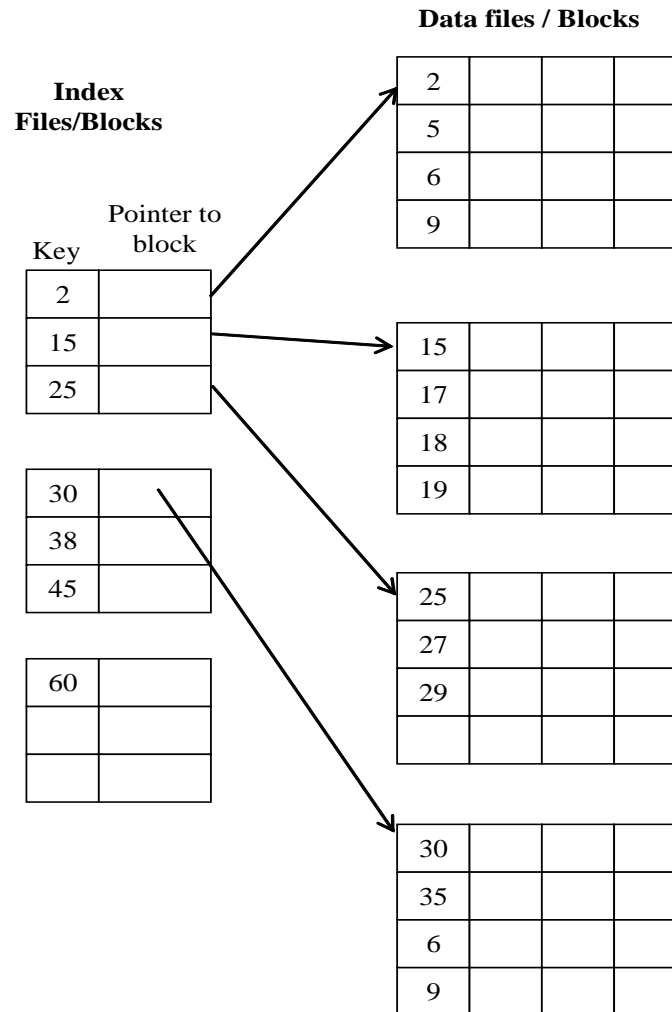
SQL Commands to create indexes:

Usually when we declare PK, an index is created automatically.

```
CREATE INDEX EMP_IND ON EMP(eid);
```

```
DROP INDEX EMP_IND;
```


Primary Indexing



Designing a Primary index



Ex 1:

Assume that we have an ordered file with 80000 records stored on disk. Block size is 512 Bytes. Record length is fixed and it is 70 Bytes. Key field(PK) length is 6 Bytes and block pointer is 4 Bytes. Assume unspanned record organization

Design a Primary index on primary key.

Solution:

Size of disk block = 512 Bytes; record length = 70 Bytes

Block pointer = 4 Bytes. Key field = 6 bytes; total records = 80000

No. records per block (Bfr) = $\text{floor}(512/70) = 7.31 = 7$

No. of data blocks needed = $\text{ceil}(80000/7) = 11429$


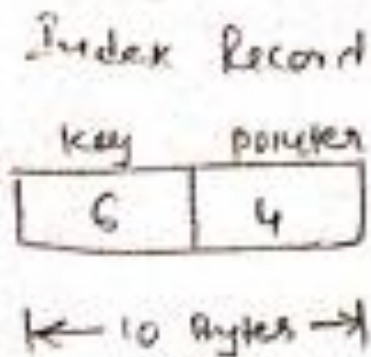
Index record length = key + pointer = $6 + 4 = 10$ Bytes

Blocking factor for index (Bfr_i) = $\text{floor}(512/10) = 51$

(known as fan-out)

No. of index blocks = $\text{Ceil}(11429/51) = 225$

No. of block accesses = $\text{ceil of } (\log_2 225) + 1 = 8 + 1 = 9$



$$\lceil \log_2 225 \rceil = 8$$

Example

Designing a Primary index



Ex 2:

Assume that we have an ordered file with 130000 records stored on disk. Block size is 1024 Bytes. Record length is fixed and it is 250 Bytes. Key field(PK) length is 10 Bytes and block pointer is 8 Bytes. Assume unspanned record organization

Design a Primary index on primary key.

Solution:

Size of disk block = 1024 Bytes; record length = 250 Bytes

Block pointer = 8 Bytes. Key field = 10 bytes; total records = 130000

No. records per block (Bfr) = $\text{floor}(1024/250) = 4.096 = 4$

No. of data blocks needed = $\text{ceil}(130000/4) = 32500$

Index record length = key + pointer = $10 + 8 = 18$ Bytes

Blocking factor for index (Bfr_i) = $\text{floor}(1024/18) = 56.88 = 56$

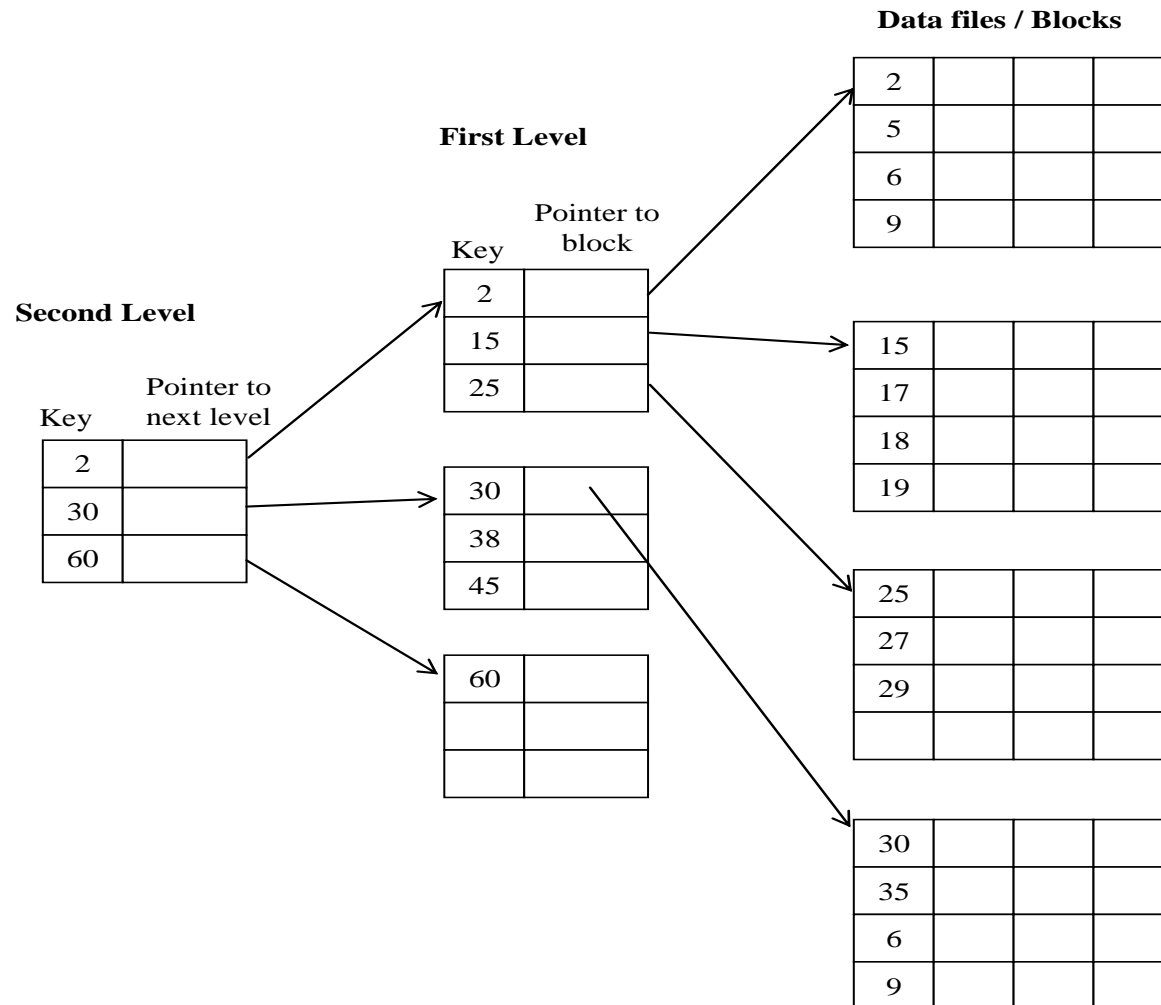
(known as fan-out)

No. of index blocks = $\text{Ceil}(32500/56) = 580.35 = 581$

No. of block accesses = $\text{ceil of } (\log_2 581) + 1 = 10 + 1 = 11$

No. of block accesses without indexing = $\text{ceil of } (\log_2 32500)$
= 15

Multilevel Indexing (Two levels)



Designing a multilevel index



Ex 2:

Assume that we have an ordered file with 80000 records stored on disk. Block size is 512 Bytes. Record length is fixed and it is 70 Bytes. Key field(PK) length is 6 Bytes and block pointer is 4 Bytes. Assume unspanned record organization

Design a multilevel index on primary key.

How many levels are there.

How many blocks are there in each index level.

Solution :

Size of the disk block=512 Bytes; record length=70 Bytes

Block pointer=4 Bytes. Key field=6 bytes; total records=80000

No. records per block(Bfr)= $\text{floor}(512/70)=7.31=7$

No. of data blocks needed= $\text{ceil}(80000/7)=11429$

Index record length= key + pointer=6+4=10 Bytes

Blocking factor for index = $\text{floor}(512/10)=51$ - fanout

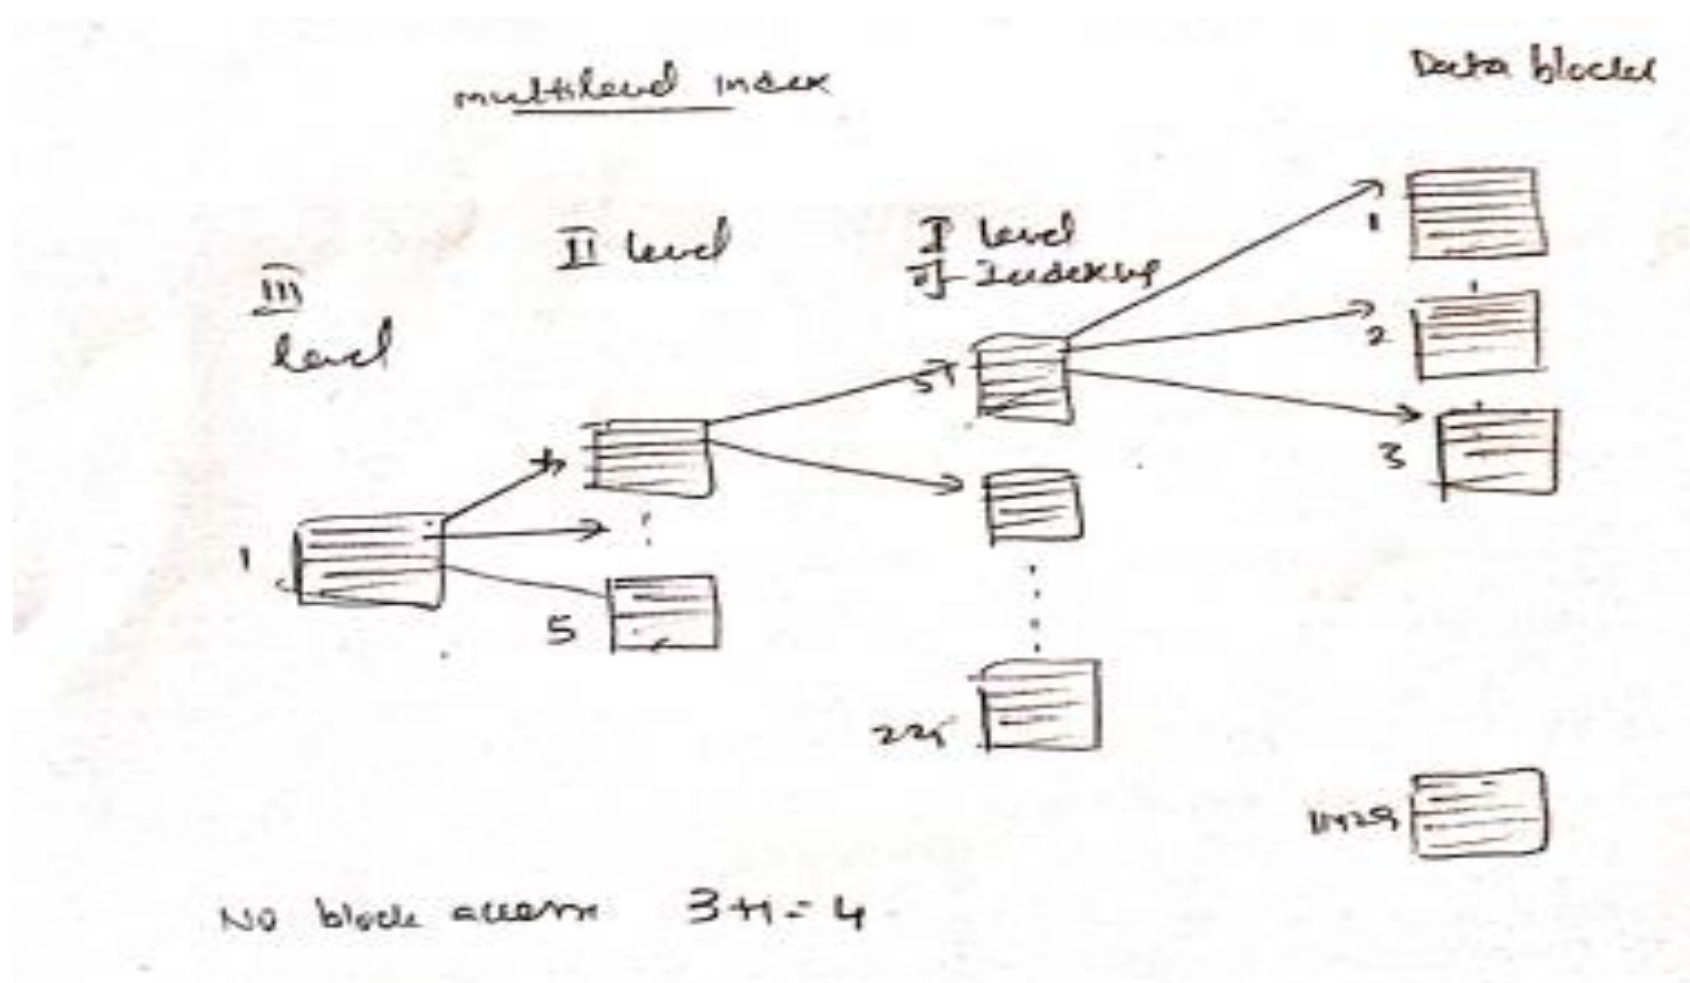
No. of index blocks in first level= $\text{Ceil}(11429/51)=225$

No. of index blocks in 2nd level= $\text{Ceil}(225/51)=5$

No. of index blocks in 3rd level= $\text{Ceil}(5/51)=1$ top level

No. of levels in indexing structure= $t=3$

No. of block accesses= No. index levels + 1= $t+1=4$



Example

Designing a multilevel index



Ex 2:

Assume that we have an ordered file with 130000 records stored on disk. Block size is 1024 Bytes. Record length is fixed and it is 250 Bytes. Key field(PK) length is 10 Bytes and block pointer is 8 Bytes. Assume unspanned record organization

Design a Primary index on primary key.

Solution:

Size of disk block = 1024 Bytes; record length = 250 Bytes

Block pointer = 8 Bytes. Key field = 10 bytes; total records = 130000

No. records per block (Bfr) = $\text{floor}(1024/250) = 4.096 = 4$

No. of data blocks needed = $\text{ceil}(130000/4) = 32500$

Index record length = key + pointer = $10 + 8 = 18$ Bytes

Blocking factor for index (Bfr_i) = $\text{floor}(1024/18) = 56.88 = 56$

(known as fan-out)

No. of index blocks (1st level) = $\text{Ceil}(32500/56) = 580.35 = 581$

No. of index blocks (2nd level) = $\text{Ceil}(581/56) = 10.375 = 11$

No. of index blocks (3rd level) = $\text{Ceil}(11/56) = 1$

No. of block accesses = No of levels + 1 = $3 + 1 = 4$

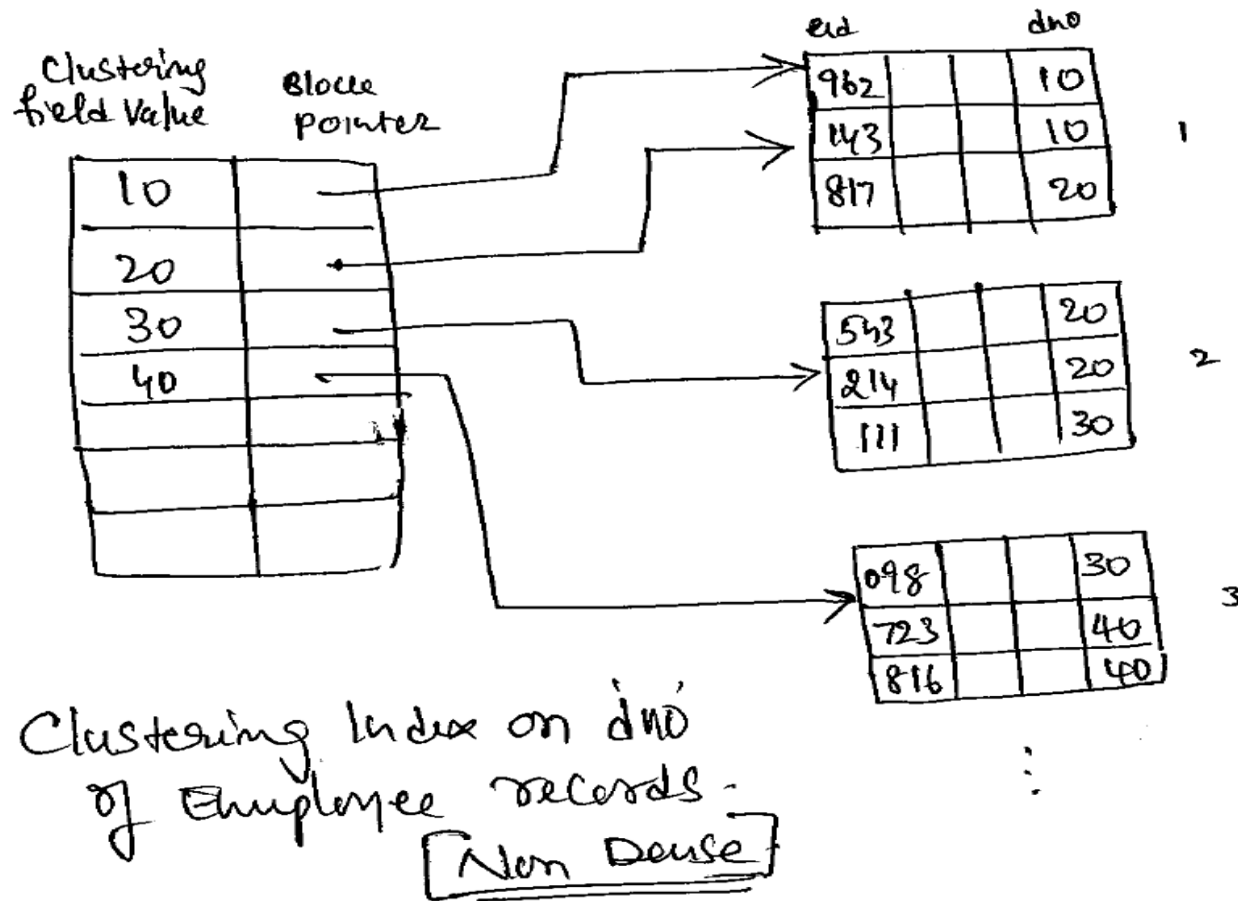
Note:

No. of block accesses without indexing = $\text{ceil}(\log_2 32500) = 15$

No. of block accesses with only one level = $\text{ceil}(\log_2 581) + 1 = 10 + 1 = 11$

No. of block accesses with only two levels = $\text{ceil}(\log_2 11) + 1 = 4 + 1 + 1 = 6$

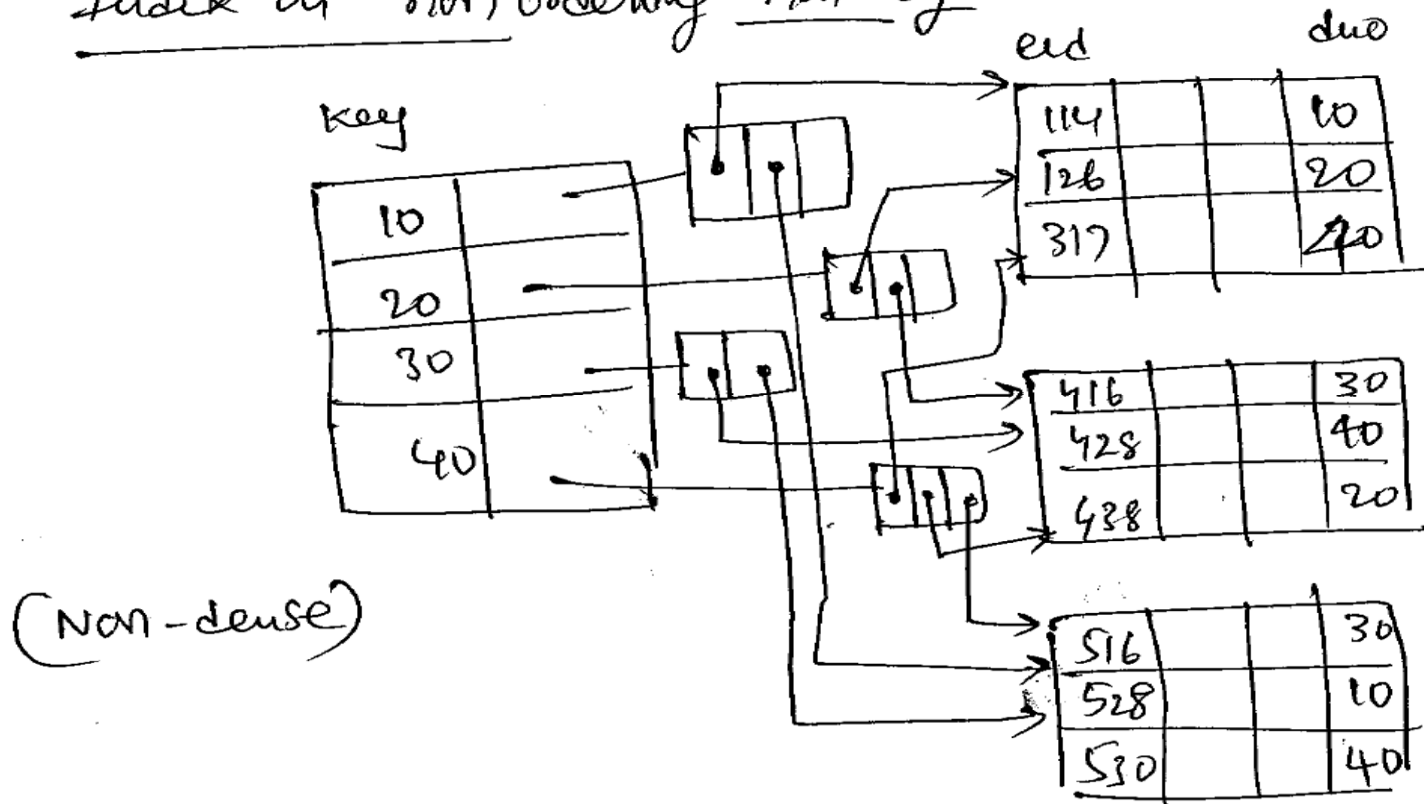
Clustering Index on Ordering Non-key



Secondary Index on Non-ordering, Non-key




Secondary Index on non ordering non-key



Secondary Index on Non-ordering key

Secondary Index on non-ordering key
(SSN) ptr



[Dense Index]

B+ Tree Indexing




B+ Tree is a multilevel search tree used to implement dynamic multilevel indexing. The primary disadvantage of implementing multilevel indexes is that the performance degrades as the file grows. It can be remedied by reorganization, but frequent reorganization is not advisable. B+ tree is best suited for multilevel indexing of files, because it is dynamic.

B+ Tree of Order p

It is a balanced tree, (all leaves are at same level).

Each internal node is of the form-



Constructing a B+ Tree

Ex 3:

Construct a B+ tree with given specifications. The order of the tree, $p=3$ and $p_{\text{leaf}}=2$. The tree should be such that all the keys in the subtree pointed by a pointer which is preceding the key must be equal to or less than the key value, and all the keys in the subtree pointed by a pointer which is succeeding the key must be greater than the key.


Insert the following keys in same order- 56, 22, 78, 42, 102, 90, 35. Show how the tree will expand after each insertion, and the final tree.

Next, delete 56, 46, 22 in the same order and show the status of the tree after each deletion.


B+ Tree P = 3 Pleaf = 2

Keys - 56, 22, 78, 42, 102, 90, 35


Insert 56, 22




Insert 78




Insert 42




Insert 102




Insert 90



Insert 35



Delete 56




Delete 46

not there
in the tree

Delete

Final tree



final tree

Node design for B+ tree

Example:

We need to design a B+ tree indexing for Student relation, on student_id attribute; the key of the relation. The attribute student_id is of 4 bytes length. Other attributes are- student_age(4 bytes), student_name(20 bytes), student_address(40 bytes), student_branch(3 bytes). The Disk block size is 1024 Bytes. If the tree-pointer takes 4 bytes, for the above situation, design the best possible number of pointers per node(internal) of the above B+ tree. Each internal node is a disk block which contains search key values and pointers to subtrees.

Solution:

Disk block size=1024 Bytes

Size of B+ tree node= size of disk block

Each tree pointer points to disk block and takes 4 Bytes.

Each key (student_id) takes 4 Bytes

In a B+ tree node, No. of pointers = no. keys +1

Assume that no. keys = n

Then no. pointers= n+1

Then min. size for a node= {(no.Keys* size of each key)+
(no.pointers * size of each pointer)} <= 1024

$$(n*4)+(n+1)*4 \leq 1024$$

$$4n+4n+4 \leq 1024$$

$$8n+4 \leq 1024$$

$$8n \leq 1024-4=1020$$

$$n \leq 127.5 \text{ or } 127$$

hence In each internal node, no. keys=127; and no. pointers=128

Note

In a B+ tree record pointer for a record with given key can be found only at leaf node.

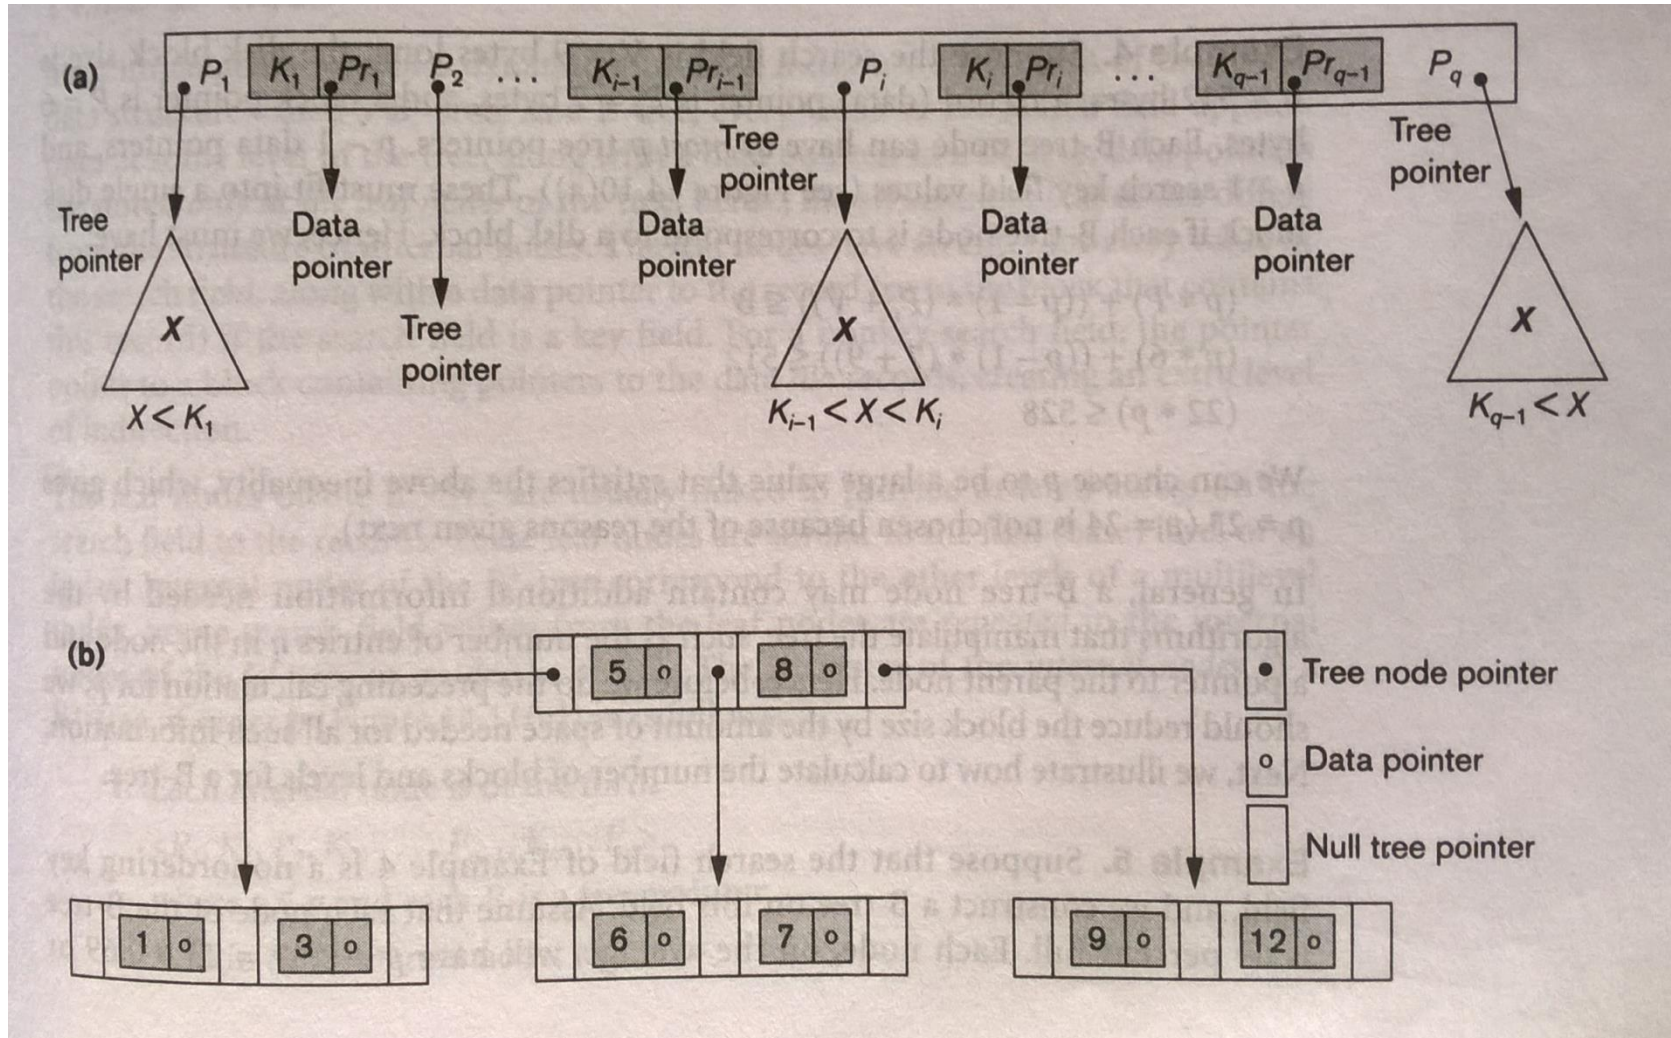
But if it is in case of B-tree it can happen at intermediate node also.

Hence in B+ tree search, success or failure can be declared only after reaching leaf_level.

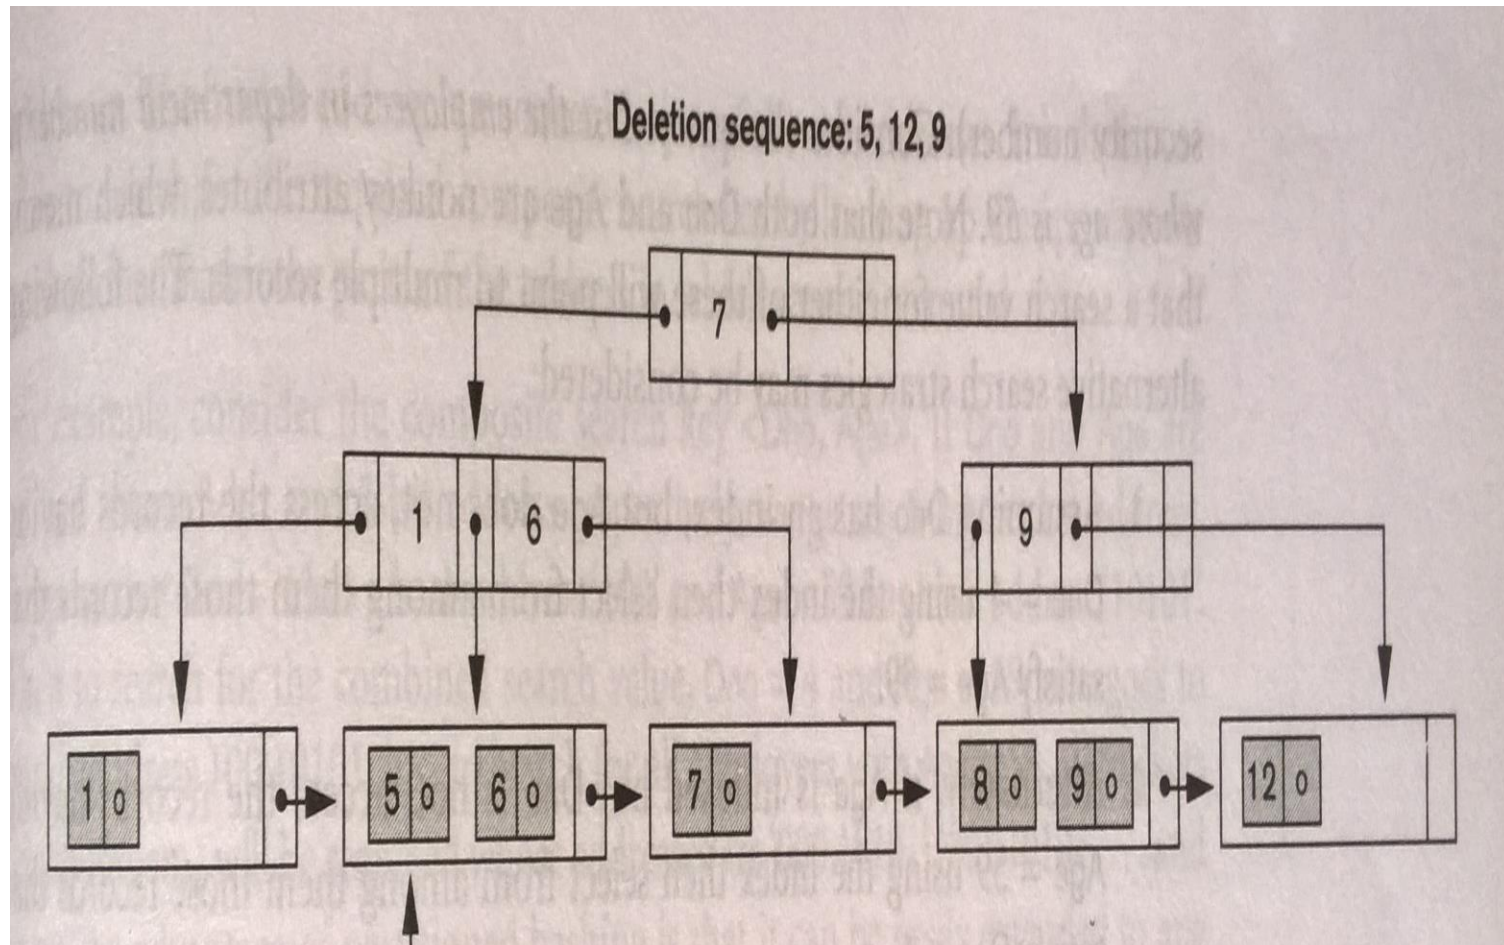
Where as in B-tree search can be successful at intermediate level as well.

On failure we reach the leaf level.

B Tree



B+ Tree



Example for B+ tree : Page 644 of T1

Reading assignment:

Bitmap Index: Ch.17.5.2 of T1 (page-649)

Summary

- ☐ *What is Indexing and its importance*
- ☐ *How Primary and Secondary indexes work*
- ☐ *Examples of Dense and Sparse Indexes*
- ☐ *What is Multilevel Indexing*
- ☐ *Some example problems on designing Primary and Multilevel Indexes*
- ☐ *What is Tree Indexing*
- ☐ *B tree and B+ tree concepts*
- ☐ *Constructing a B+ tree (Insert/Delete operations)*
- ☐ *Designing a B+ Tree node structure*
- ☐ *Bitmap index(reading assignment)*